


# Chapter 3: Transport Layer

*Computer Networking: A Top-Down Approach*

6<sup>th</sup> edition

Jim Kurose, Keith Ross

Addison-Wesley



# Chapter 3: Transport Layer

## Our goals:

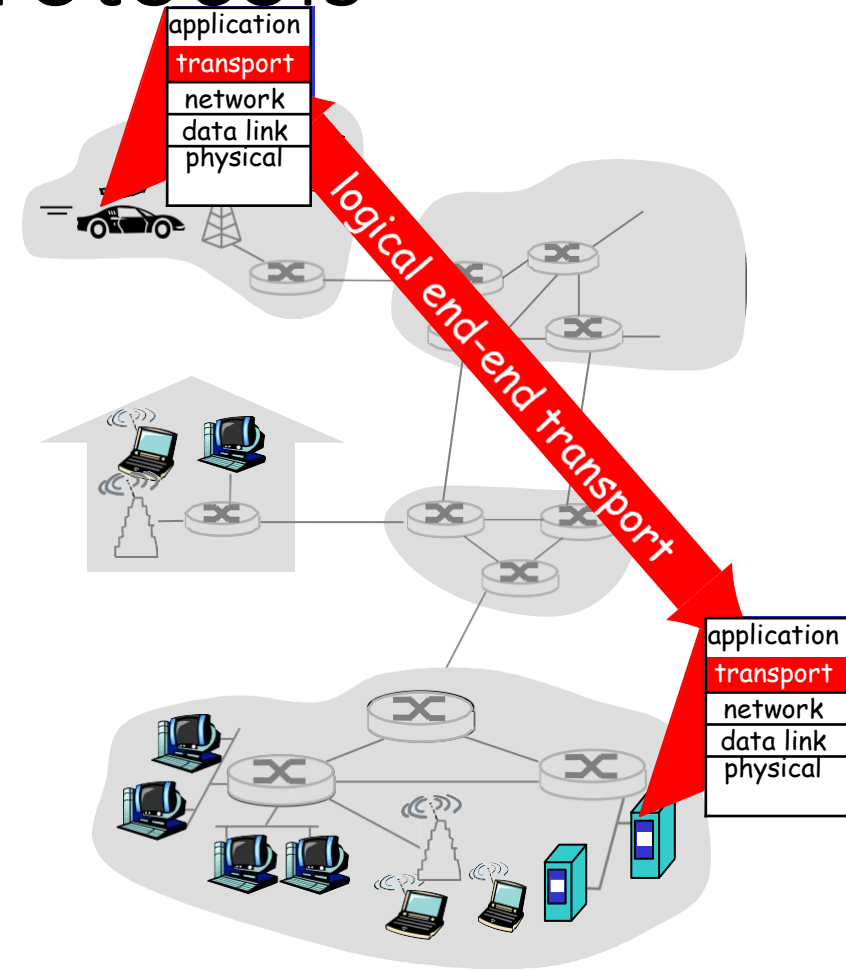
- ❖ understand principles behind transport layer services:
  - multiplexing/demultiplexing
  - reliable data transfer
  - flow control
  - congestion control
- ❖ learn about transport layer protocols in the Internet:
  - UDP: connectionless transport
  - TCP: connection-oriented transport
  - TCP congestion control

# Chapter 3 outline

1. Transport-layer services
2. Multiplexing and demultiplexing
3. Connectionless transport: UDP
4. Principles of reliable data transfer
5. Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management

# Transport services and protocols

- ❖ provide **logical communication** between app processes running on different hosts
- ❖ transport protocols run in end systems
  - send side: breaks app messages into **segments**, passes to network layer
  - rcv side: reassembles segments into messages, passes to app layer
- ❖ more than one transport protocol available to apps
  - Internet: TCP and UDP

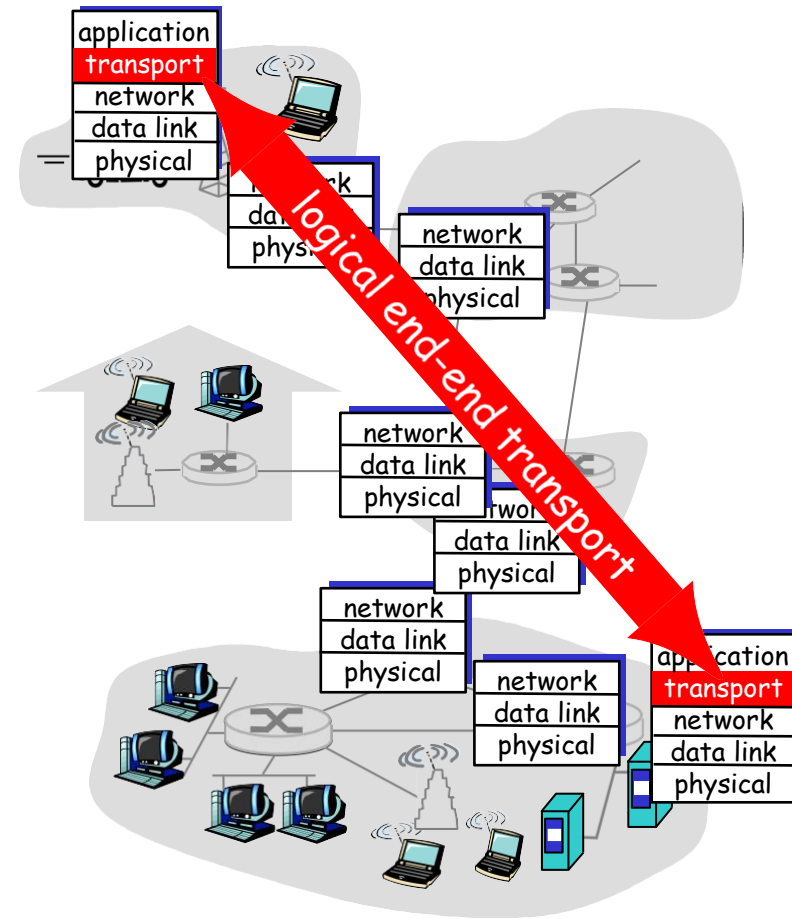


# Transport vs. network layer

- ❖ **network layer:** logical communication between hosts
- ❖ **transport layer:** logical communication between processes
  - relies on, enhances, network layer services

# Internet transport-layer protocols

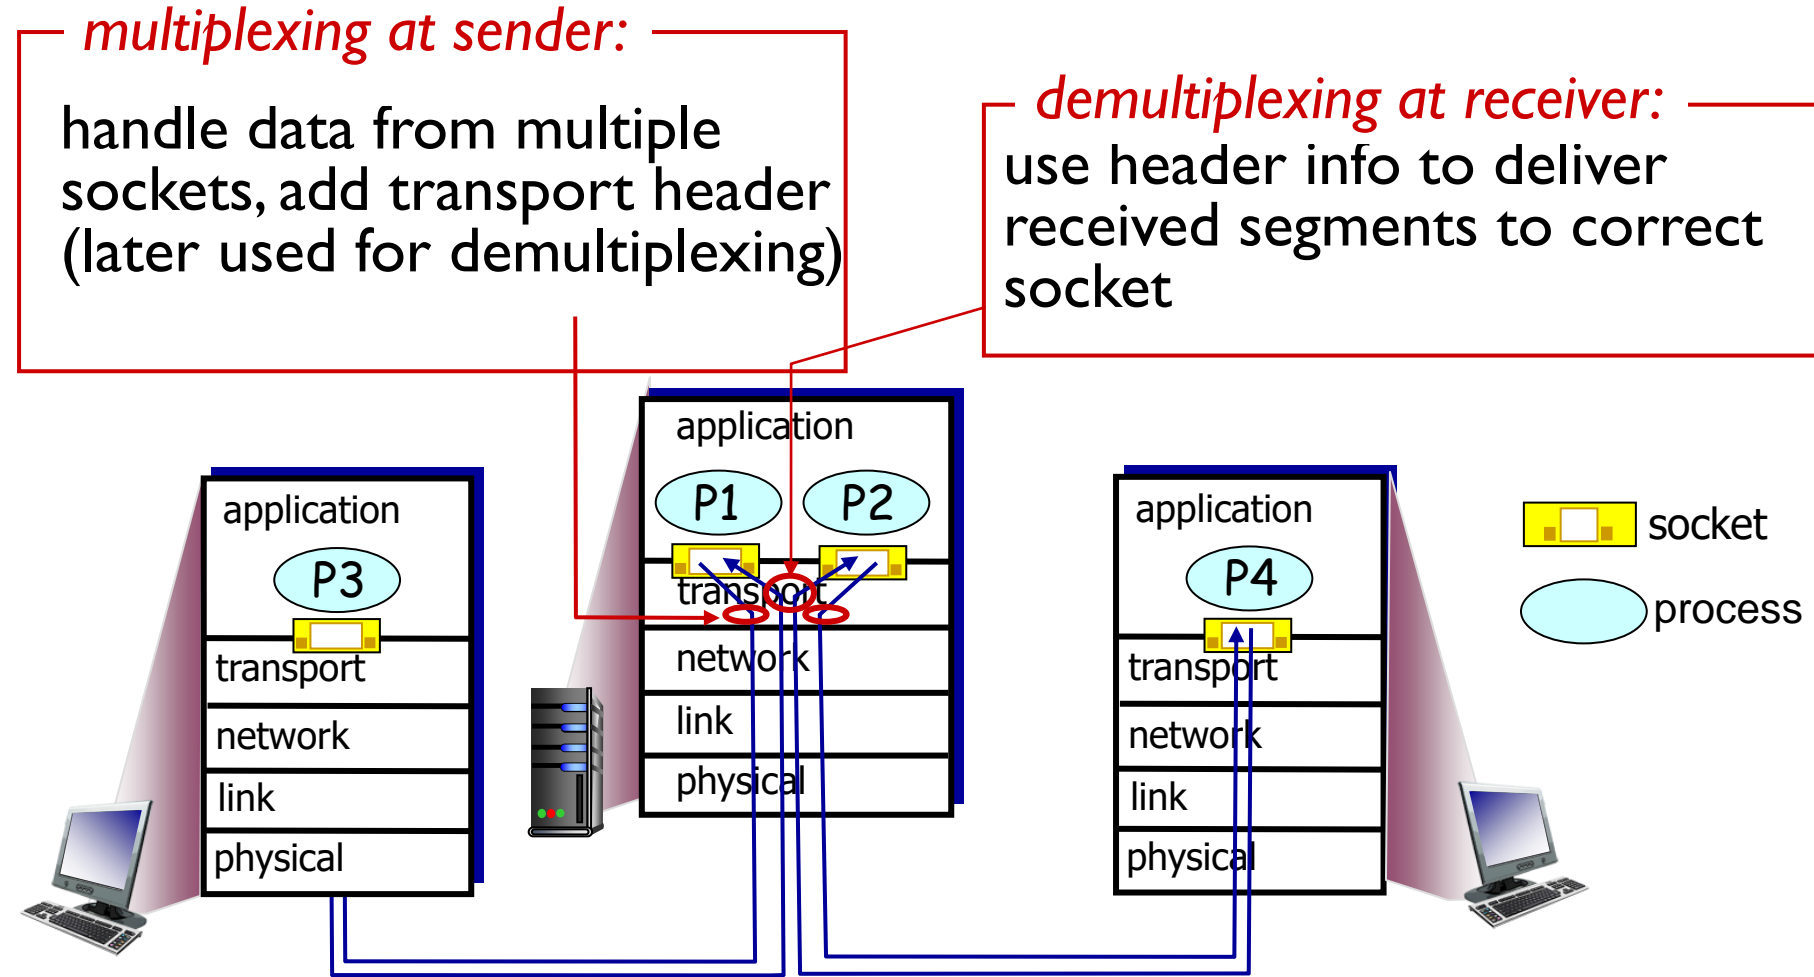
- ❖ reliable, in-order delivery (TCP)
  - congestion control
  - flow control
  - connection setup
- ❖ unreliable, unordered delivery: UDP
  - no-frills extension of "best-effort" IP
- ❖ services not available:
  - delay guarantees
  - bandwidth guarantees



# Chapter 3 outline

1. Transport-layer services
2. Multiplexing and demultiplexing
3. Connectionless transport: UDP
4. Principles of reliable data transfer
5. Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management
6. Principles of congestion control
7. TCP congestion control

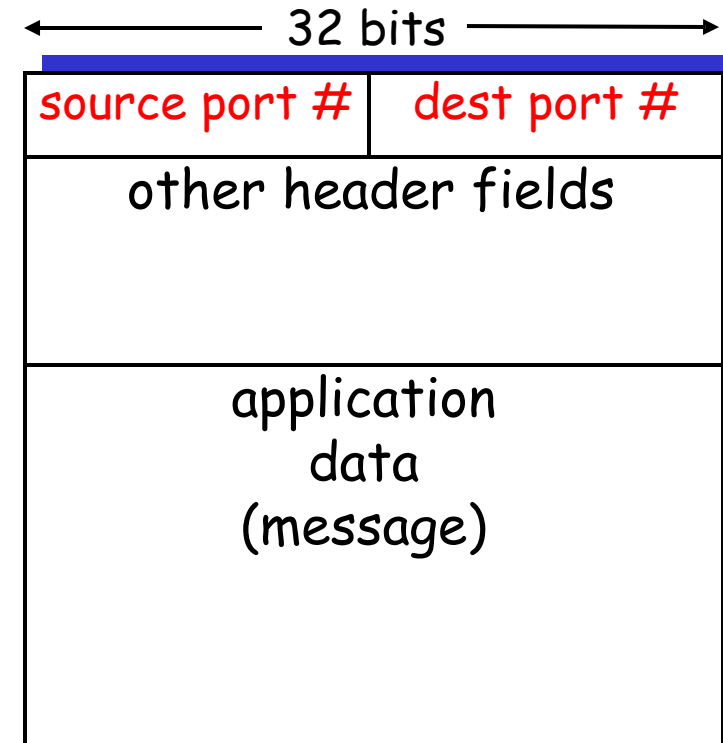
# Multiplexing/demultiplexing





# How demultiplexing works

- ❖ **host receives IP datagrams**
  - each datagram has source IP address, destination IP address
  - each datagram carries 1 transport-layer segment
  - each segment has source, destination port number
- ❖ **host uses IP addresses & port numbers to direct segment to appropriate socket**



TCP/UDP segment format

# Connectionless demultiplexing

❖ *recall*: created socket has host-local port #:

```
DatagramSocket mySocket1  
= new DatagramSocket(12534) ;
```

❖ *recall*: when creating datagram to send into UDP socket, must specify

- destination IP address
- destination port #

❖ when host receives UDP segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



IP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at dest

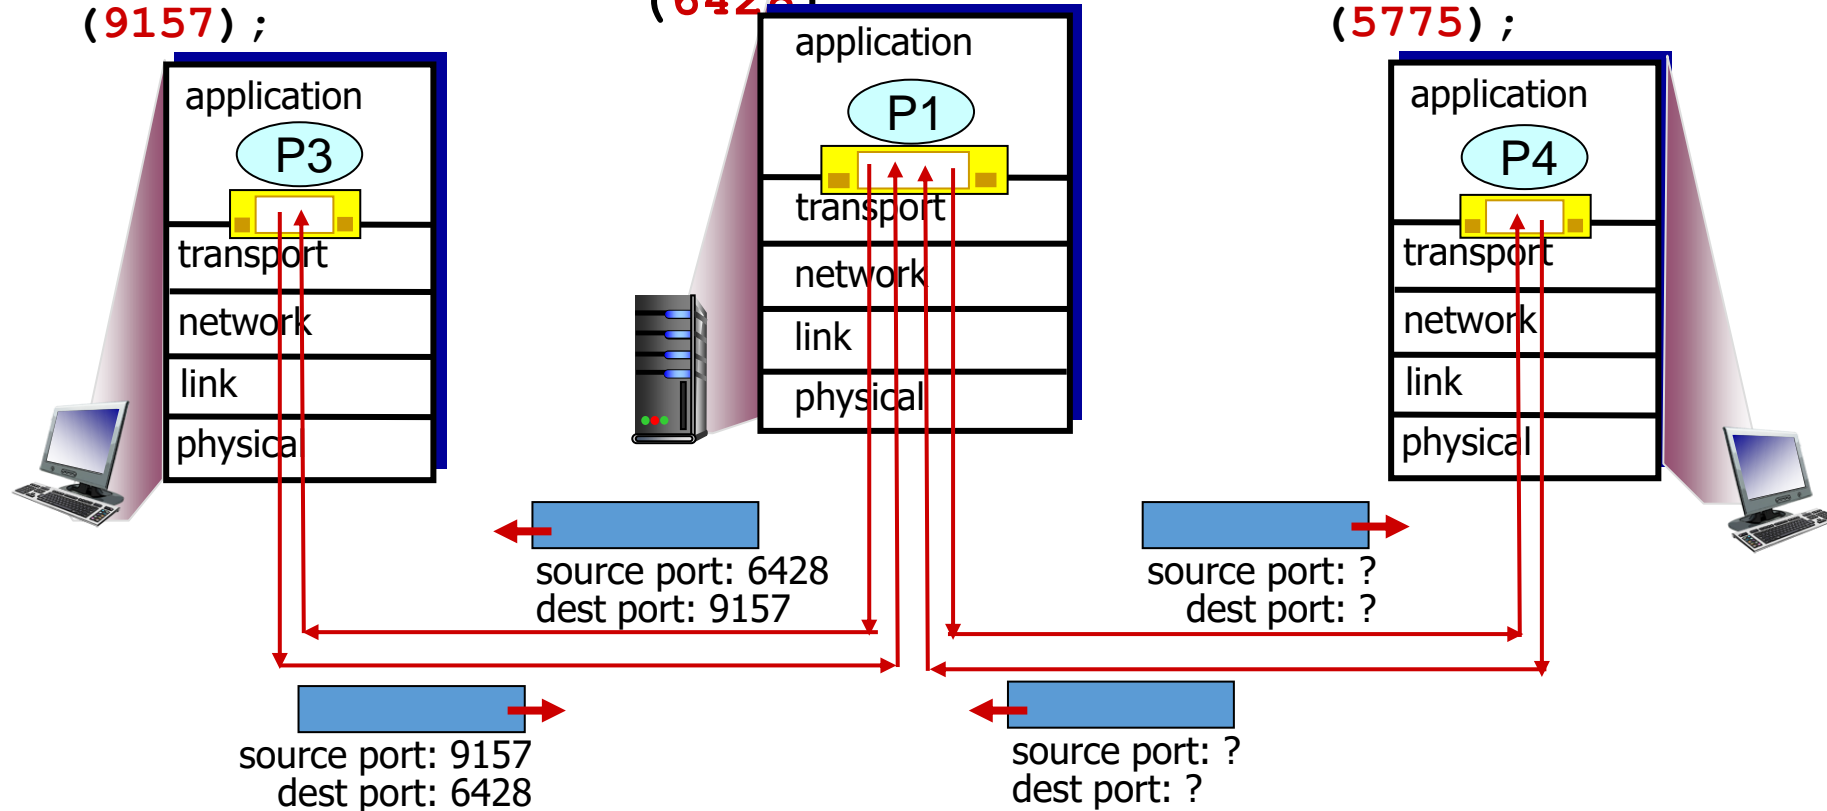
# Connectionless demux (cont)

DatagramSocket

```
DatagramSocket  
mySocket2 = new  
DatagramSocket  
(9157);
```

```
serverSocket = new  
DatagramSocket  
(6428);
```

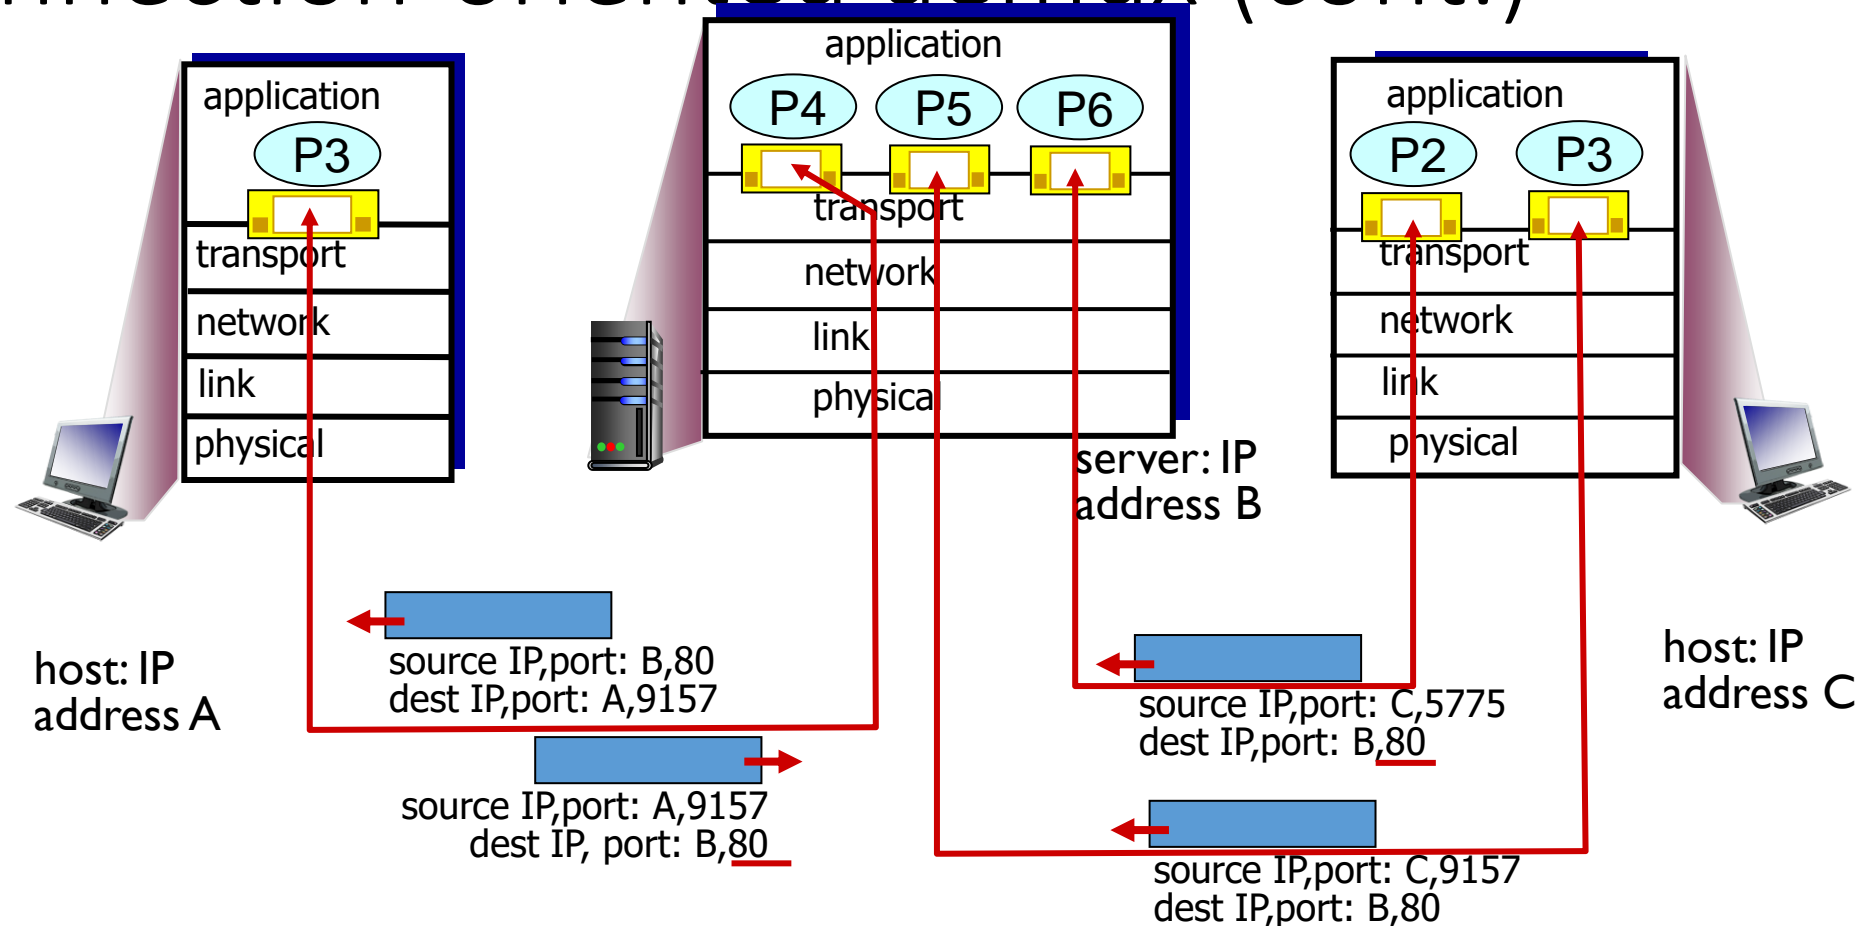
```
DatagramSocket  
mySocket1 = new  
DatagramSocket  
(5775);
```



# Connection-oriented demux

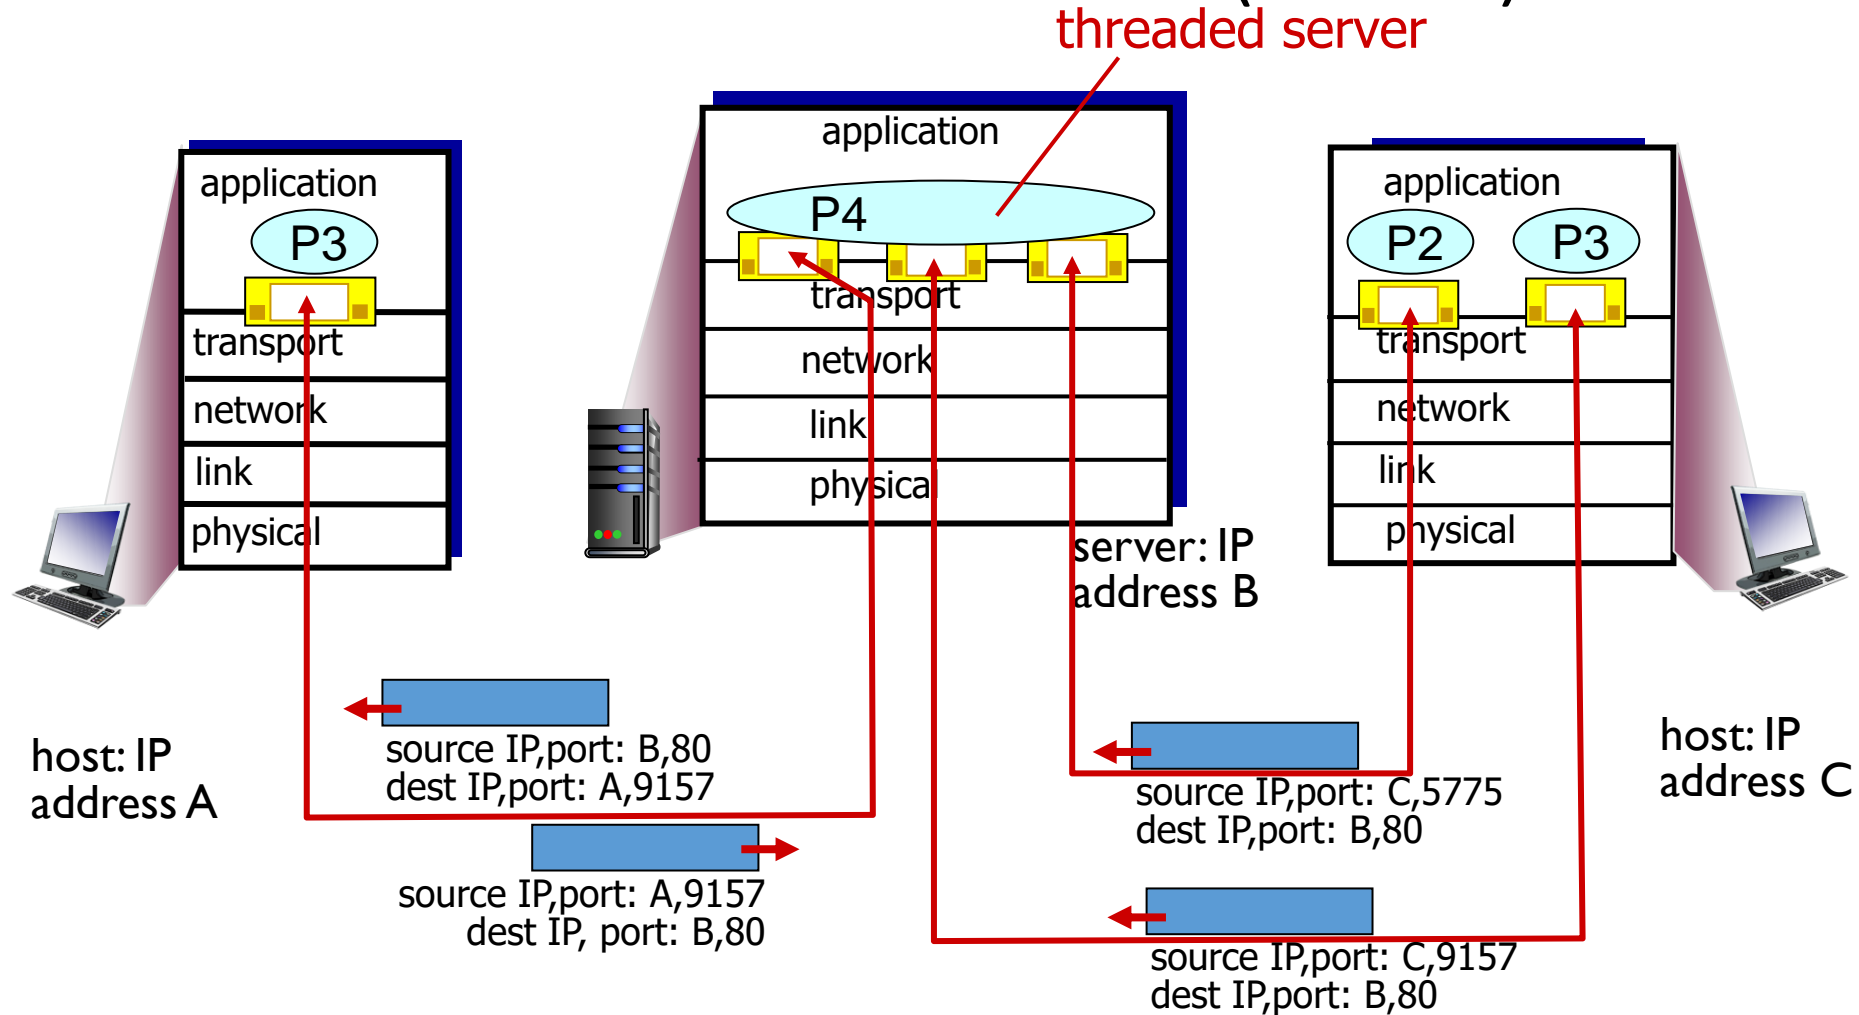
- ❖ TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- ❖ recv host uses all four values to direct segment to appropriate socket
- ❖ server host may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
- ❖ web servers have different sockets for each connecting client
  - non-persistent HTTP will have different socket for each request

# Connection-oriented demux (cont.)



three segments, all destined to IP address: B,  
dest port: 80 are demultiplexed to *different* sockets

# Connection-oriented demux (cont.)



# Chapter 3 outline

1. Transport-layer services
2. Multiplexing and demultiplexing
3. Connectionless transport: UDP
4. Principles of reliable data transfer
5. Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management

# UDP: User Datagram Protocol [RFC 768]

- ❖ “no frills,” “bare bones” Internet transport protocol
- ❖ “best effort” service, UDP segments may be:
  - lost
  - delivered out of order to app
- ❖ **connectionless:**
  - no handshaking between UDP sender, receiver
  - each UDP segment handled independently of others

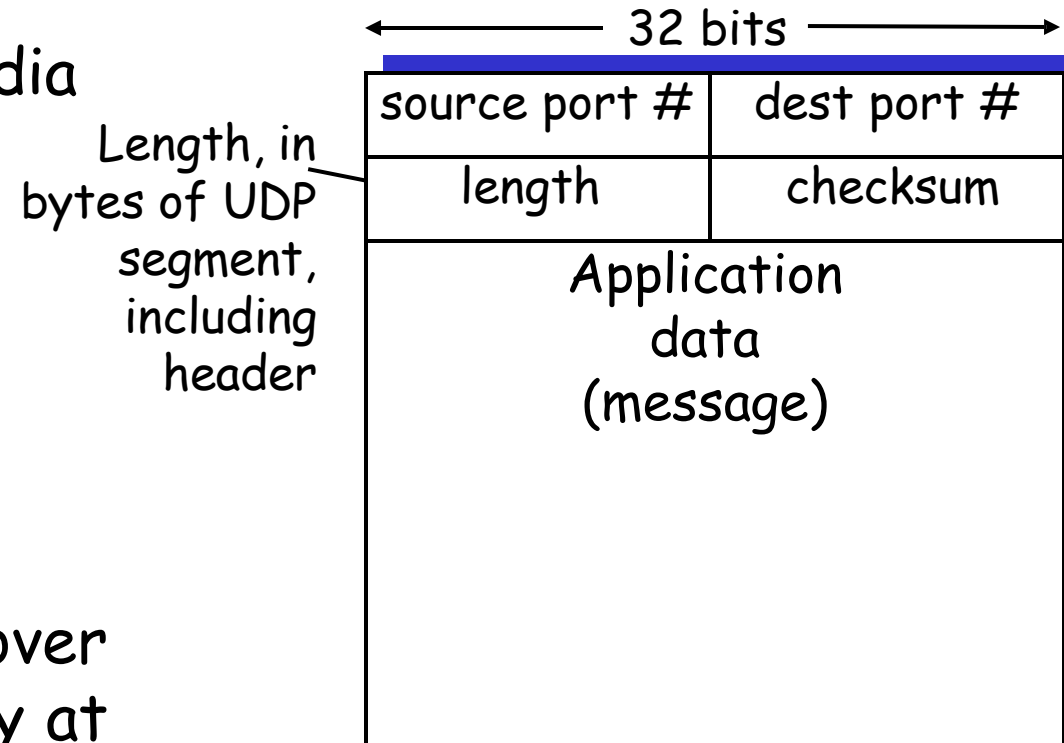
## Why is there a UDP?

- ❖ no connection establishment (which can add delay)
- ❖ simple: no connection state at sender, receiver
- ❖ small segment header
- ❖ no congestion control: UDP can blast away as fast as desired



# UDP: more

- ❖ often used for streaming multimedia apps
  - loss tolerant
  - rate sensitive
- ❖ other UDP uses
  - DNS
  - SNMP
- ❖ reliable transfer over UDP: add reliability at application layer
  - application-specific error recovery!



UDP segment format

# UDP checksum

Goal: detect "errors" (e.g., flipped bits) in transmitted segment

## Sender:

- ❖ treat segment contents as sequence of 16-bit integers
- ❖ checksum: addition (1's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

## Receiver:

- ❖ compute checksum of received segment
- ❖ check if computed checksum equals checksum field value:
  - NO - error detected
  - YES - no error detected.  
But maybe errors nonetheless? More later  
....

# Internet Checksum Example

example: add three words 16-bit

```
0110011001100110
0101010101010101
0000111100001111
```

---

```
sum      1100101011001010
checksum 0011010100110101
```


*Note:* when adding numbers, a carryout from the most significant bit needs to be added to the result

# Chapter 3 outline

1. Transport-layer services
2. Multiplexing and demultiplexing
3. Connectionless transport: UDP
4. Principles of reliable data transfer
5. Connection-oriented transport: TCP
  - segment structure
  - reliable data transfer
  - flow control
  - connection management

# Principles of Reliable data transfer

- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!

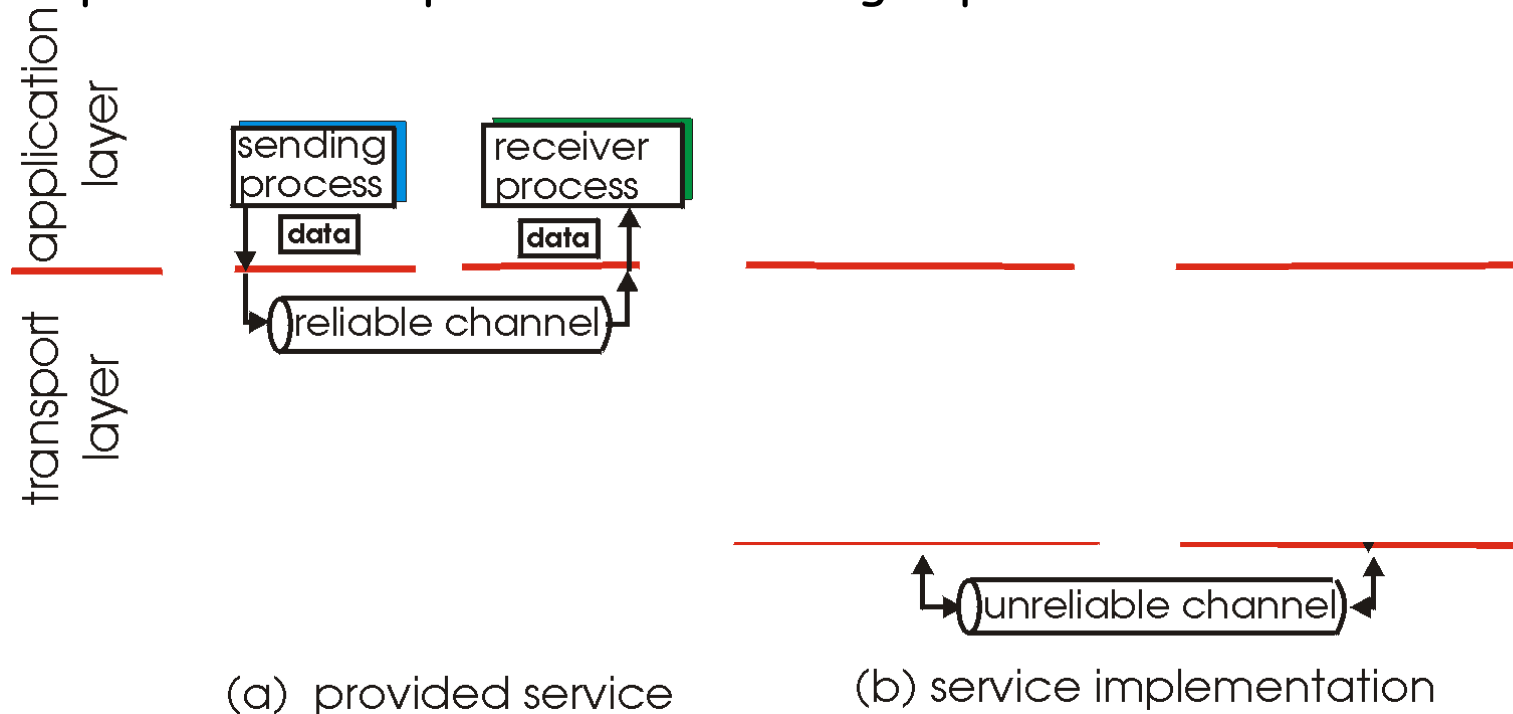


(a) provided service

- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Principles of Reliable data transfer

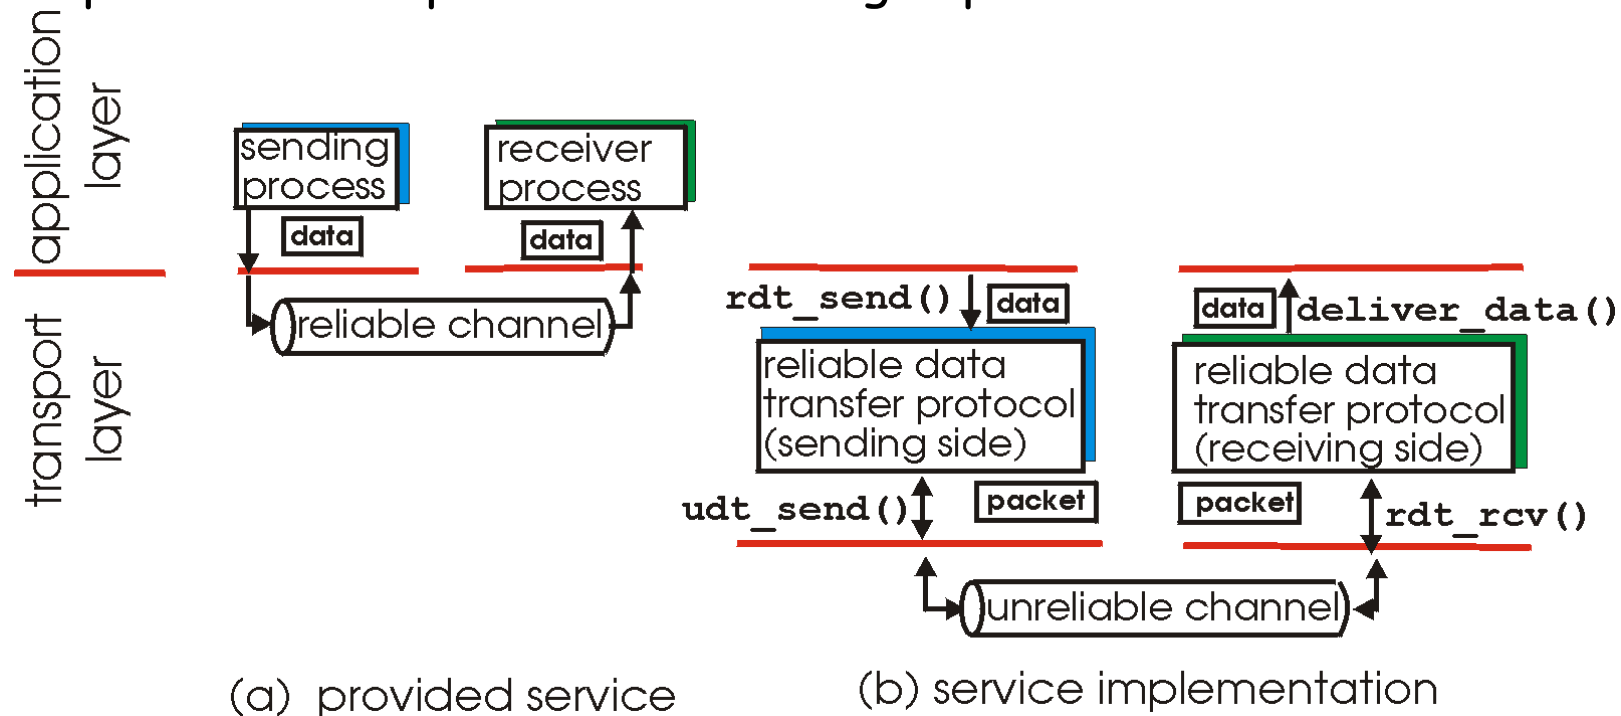
- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

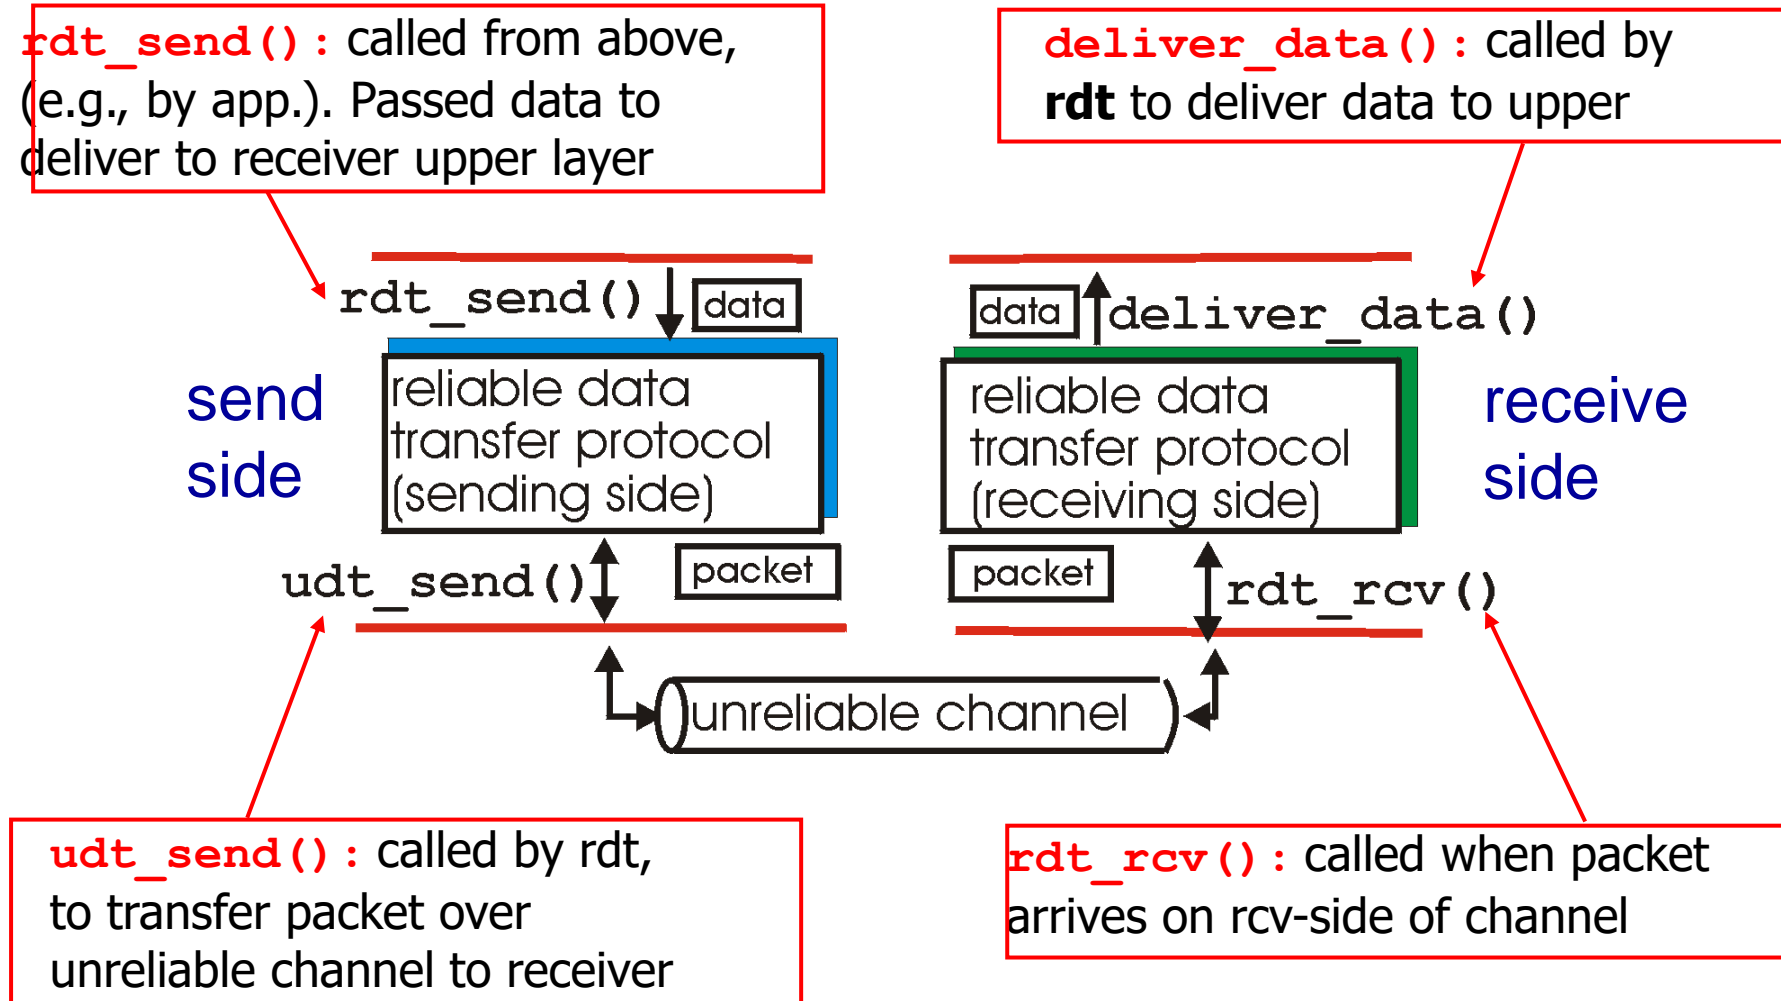
# Principles of Reliable data transfer

- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!



- ❖ characteristics of unreliable channel will determine complexity of reliable data transfer protocol (rdt)

# Reliable data transfer: getting started

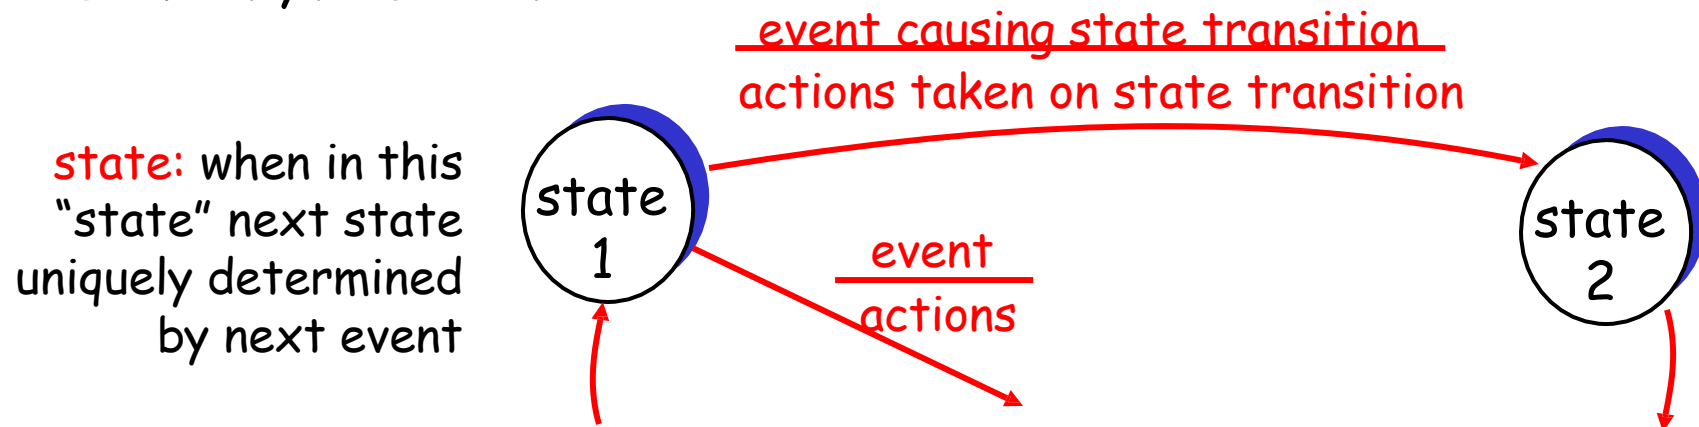




# Reliable data transfer: getting started


**We'll:**

- ❖ incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ consider only unidirectional data transfer
  - but control info will flow on both directions!
- ❖ use finite state machines (FSM) to specify sender, receiver




# Rdt1.0: reliable transfer over a reliable channel

- ❖ underlying channel perfectly reliable
  - no bit errors
  - no loss of packets
- ❖ separate FSMs for sender, receiver:
  - sender sends data into underlying channel
  - receiver read data from underlying channel



sender



receiver

# Rdt2.0: channel with bit errors


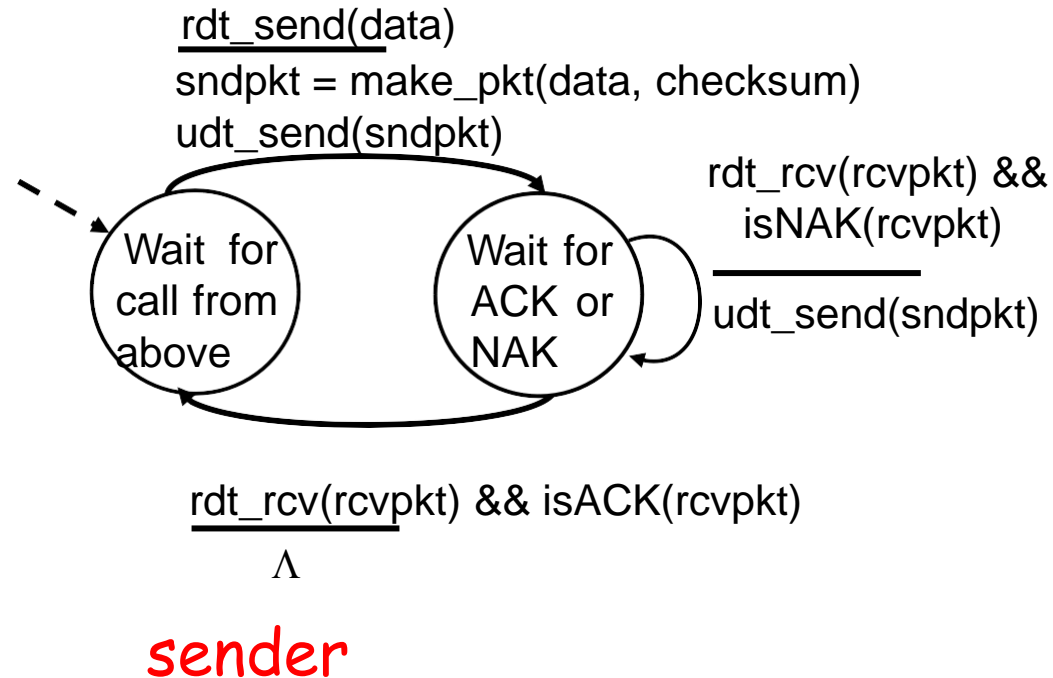
- ❖ underlying channel may flip bits in packet
  - checksum to detect bit errors
- ❖ the question: how to recover from errors:

*How do humans recover from "errors"  
during conversation?*

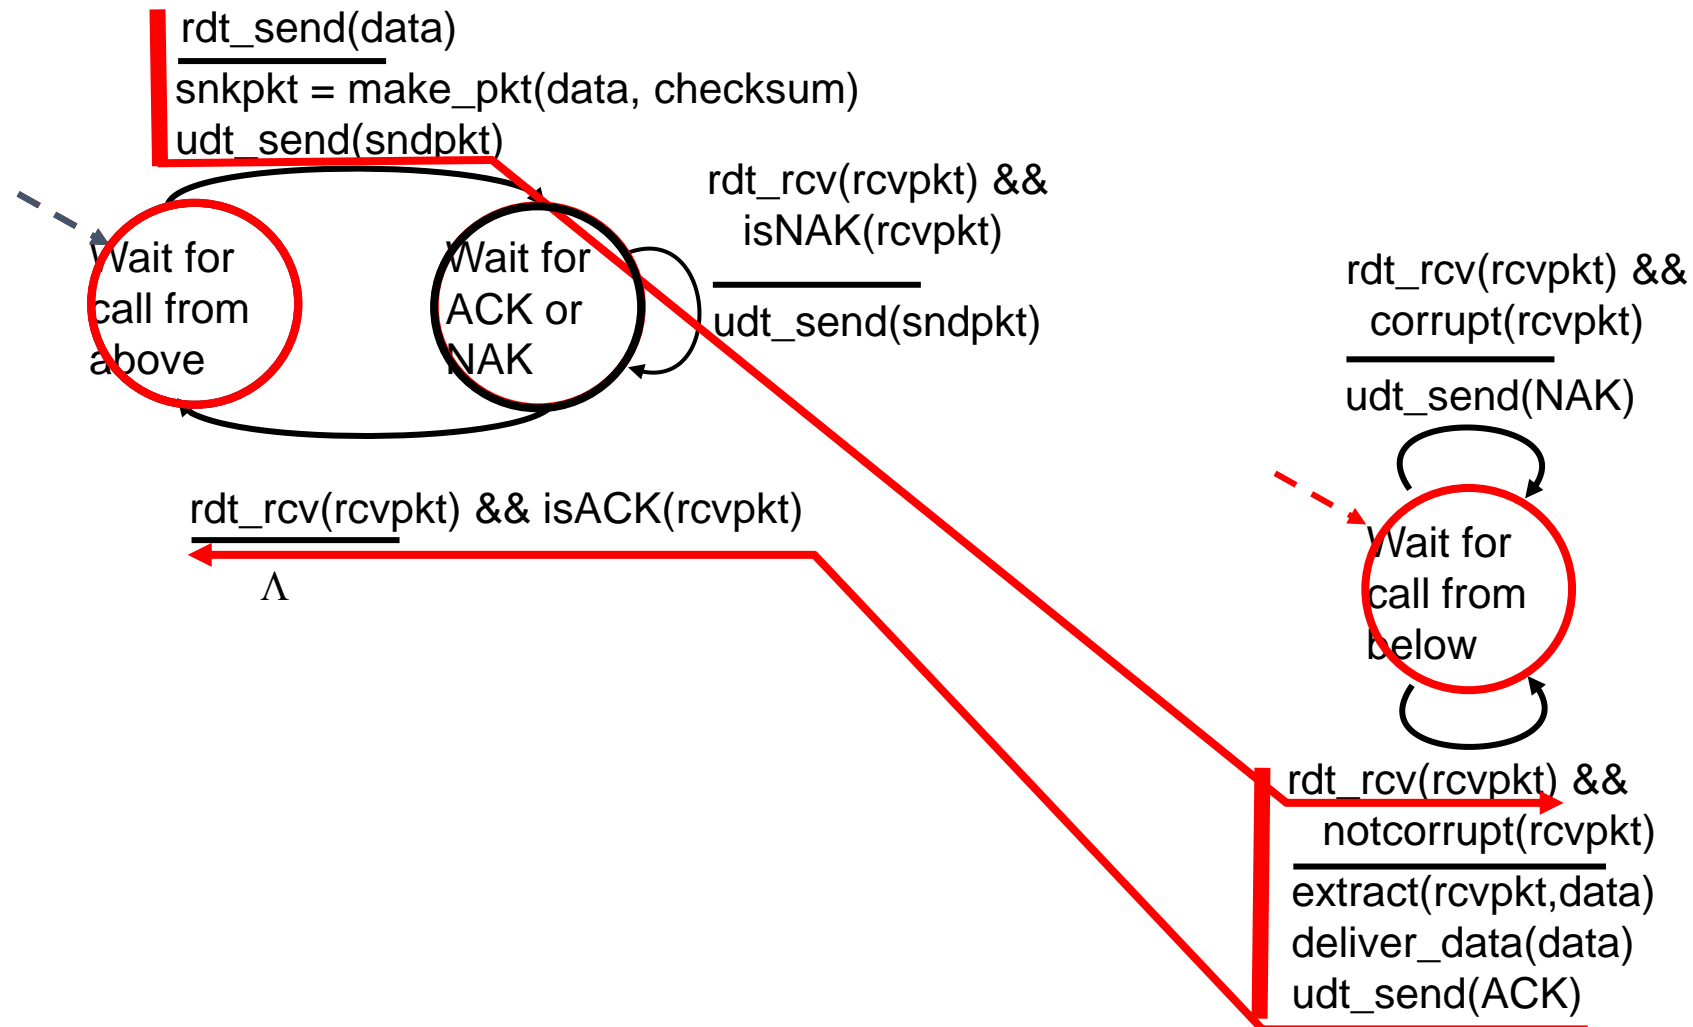
# Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
  - checksum to detect bit errors
- ❖ the question: how to recover from errors:
  - **acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
  - **negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors
  - sender retransmits pkt on receipt of NAK
- ❖ new mechanisms in `rdt2.0` (beyond `rdt1.0`):
  - error detection (checksum)
  - receiver feedback: control msgs (ACK,NAK) rcvr->sender

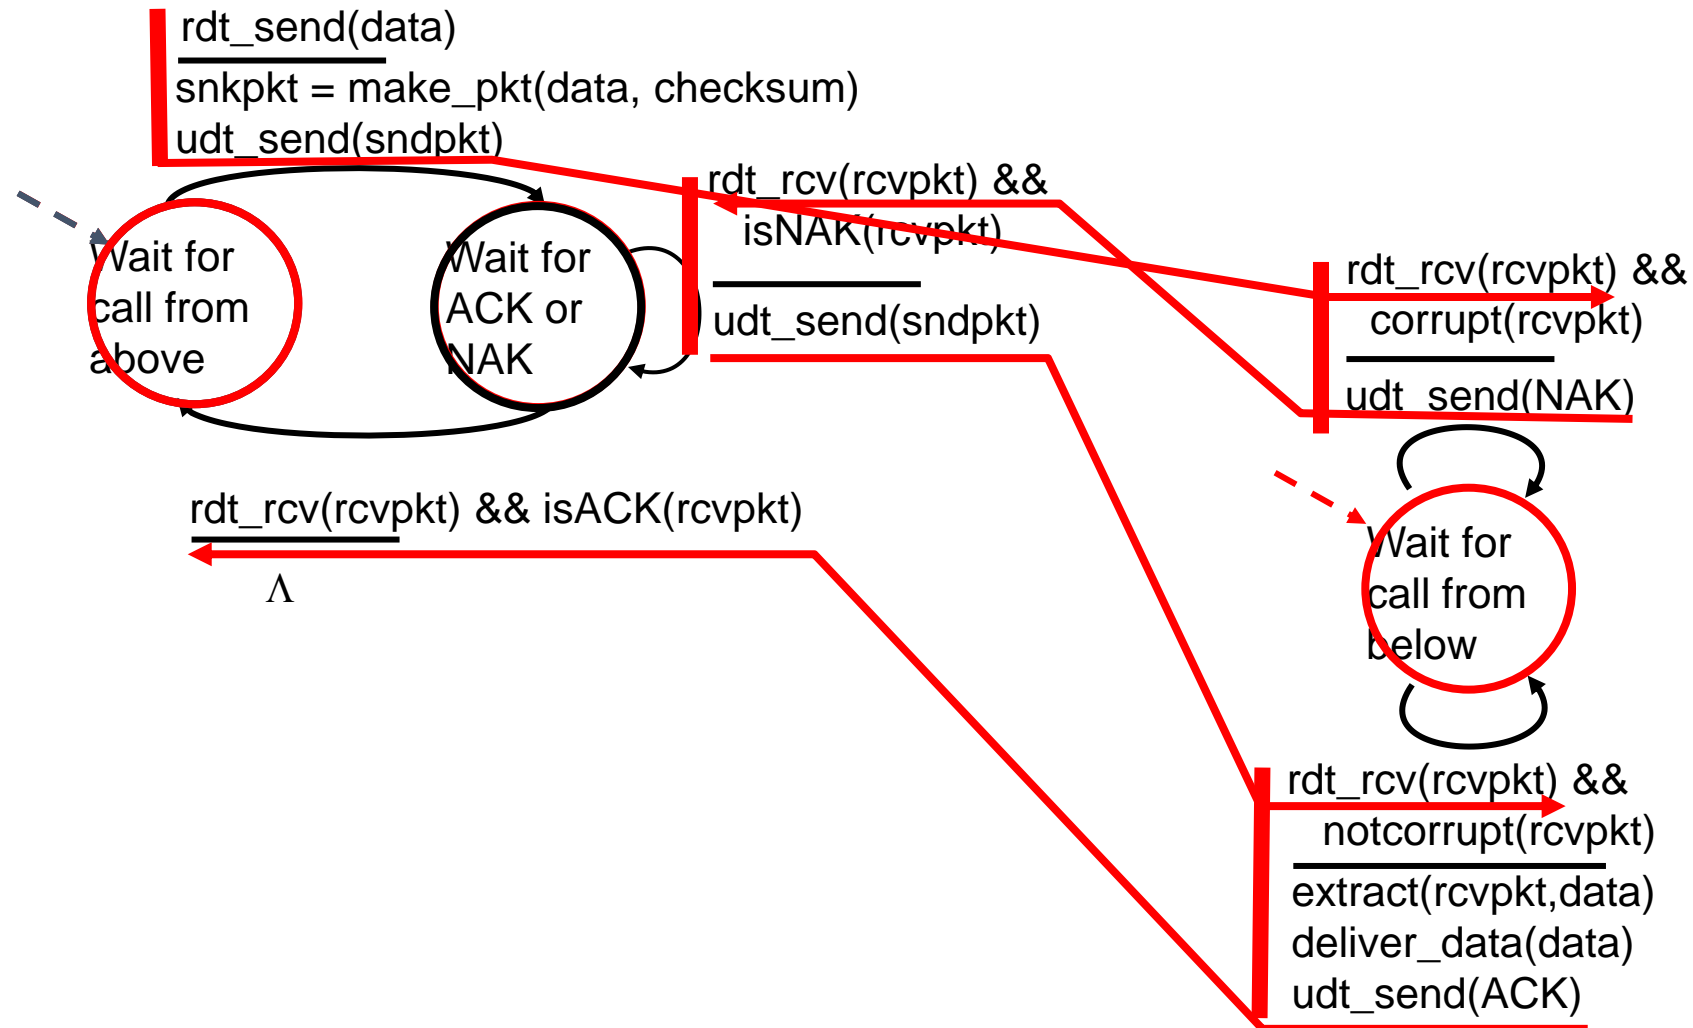
# rdt2.0: FSM specification



# rdt2.0: operation with no errors



# rdt2.0: error scenario



# rdt2.0 has a fatal flaw!

## what happens if ACK/NAK corrupted?

- ❖ sender doesn't know what happened at receiver!
- ❖ can't just retransmit: possible duplicate

## handling duplicates:

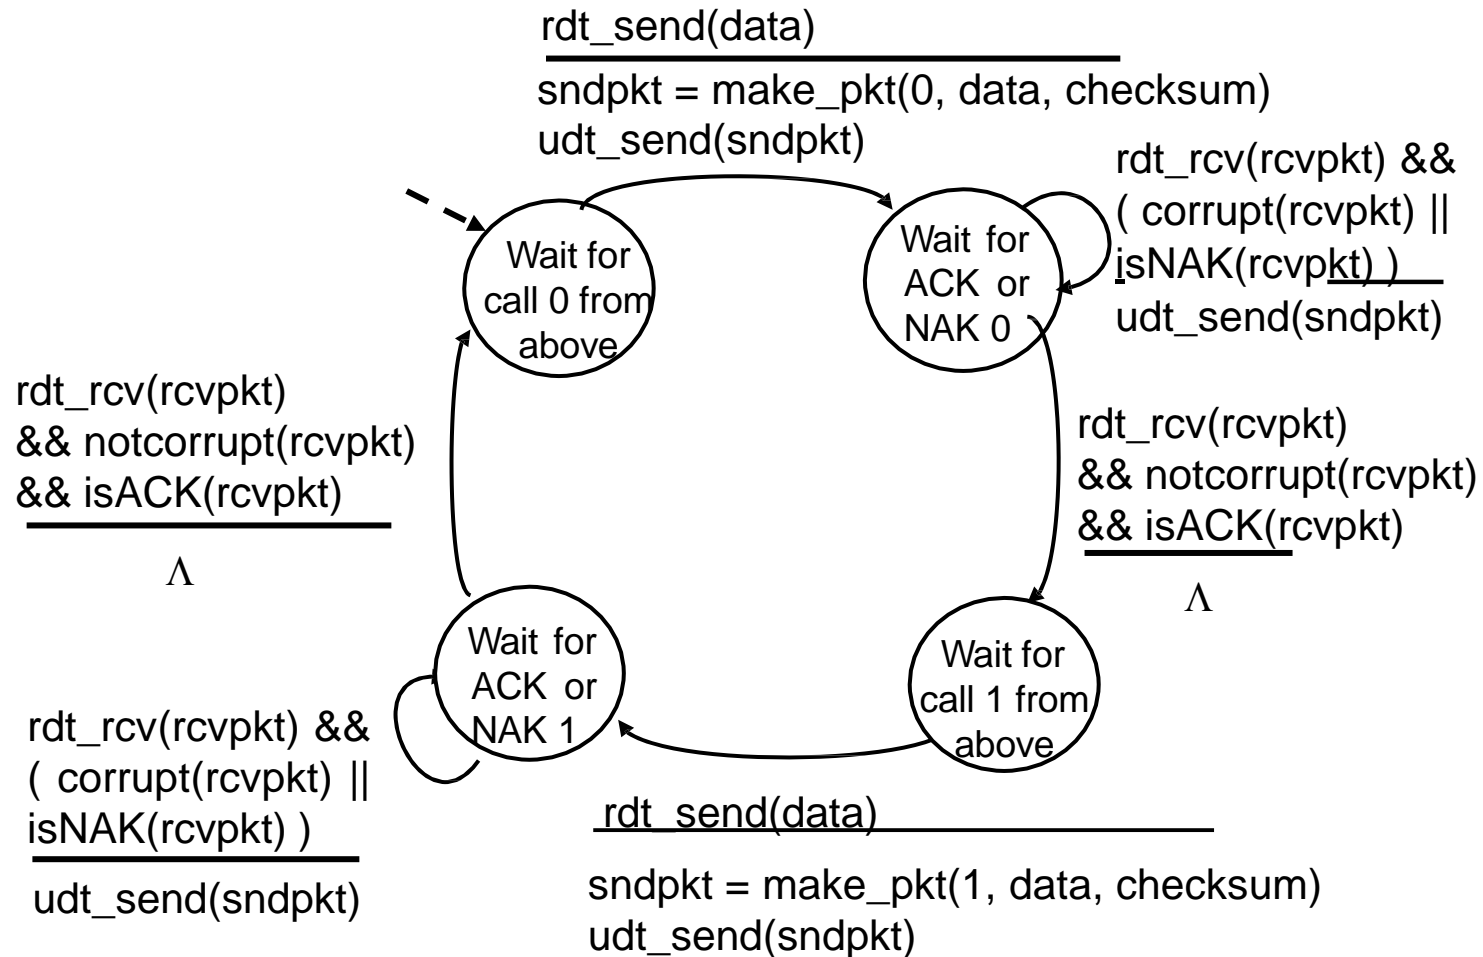
- ❖ sender retransmits current pkt if ACK/NAK corrupted
- ❖ sender adds *sequence number* to each pkt
- ❖ receiver discards (doesn't deliver up) duplicate pkt

## stop and wait

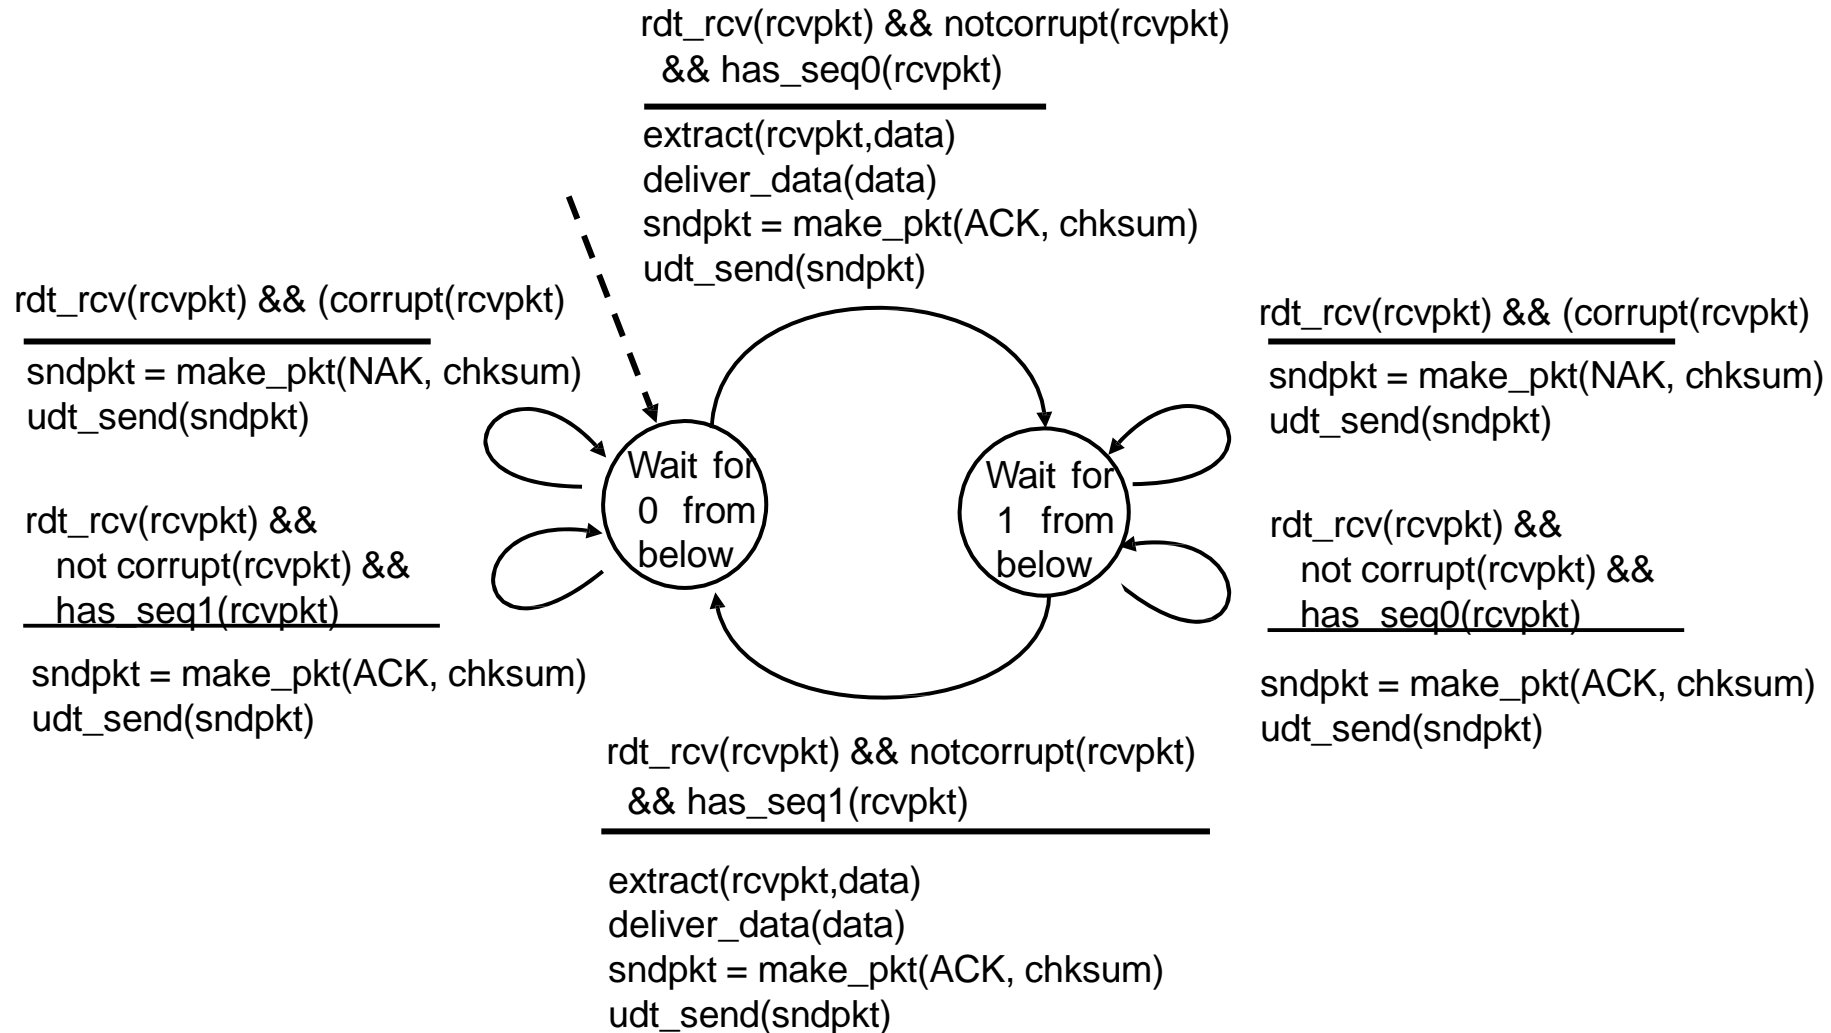
sender sends one packet,  
then waits for receiver  
response



# rdt2.1: sender, handles garbled ACK/NAKs



# rdt2.1: receiver, handles garbled ACK/NAKs



# rdt2.1: discussion

## Sender:

- ❖ seq # added to pkt
- ❖ two seq. #'s (0,1) will suffice.
- ❖ must check if received ACK/NAK corrupted
- ❖ twice as many states
  - state must "remember" whether "current" pkt has 0 or 1 seq. #

## Receiver:

- ❖ must check if received packet is duplicate
  - state indicates whether 0 or 1 is expected pkt seq #
- ❖ note: receiver can not know if its last ACK/NAK received OK at sender

# Recap

Rdt Version	Scenario	Features Added
1.0	No error	Nothing
2.0	Data Bit Error	Checksum, ACK/NAK
2.1	Data Bit Error ACK/NAK Bit Error	Checksum, ACK/NAK, Sequence Number
3.0	What else?	

# rdt3.0: channels with errors and loss

## New assumption:

underlying channel can also lose packets (data or ACKs)

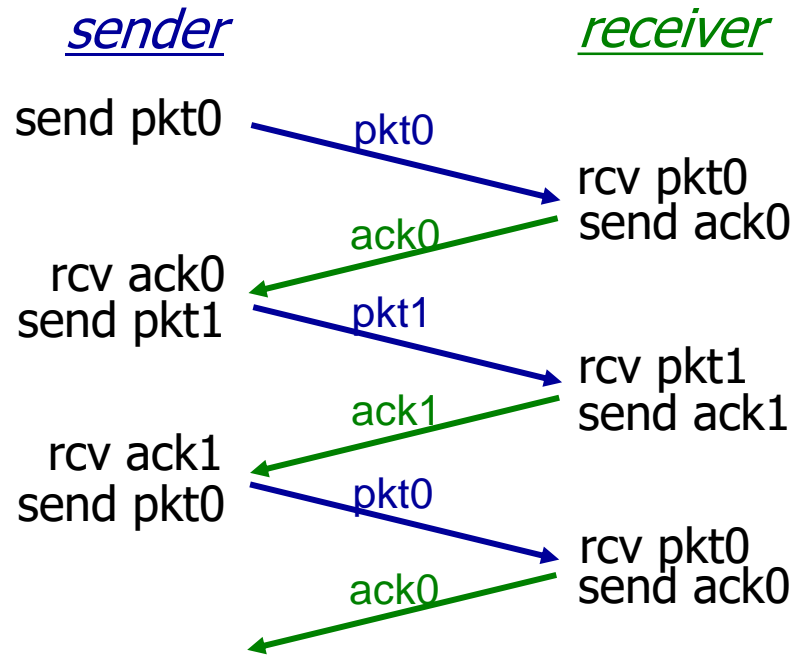
- checksum, seq. #, ACKs, retransmissions will be of help, but not enough

## Approach: sender waits

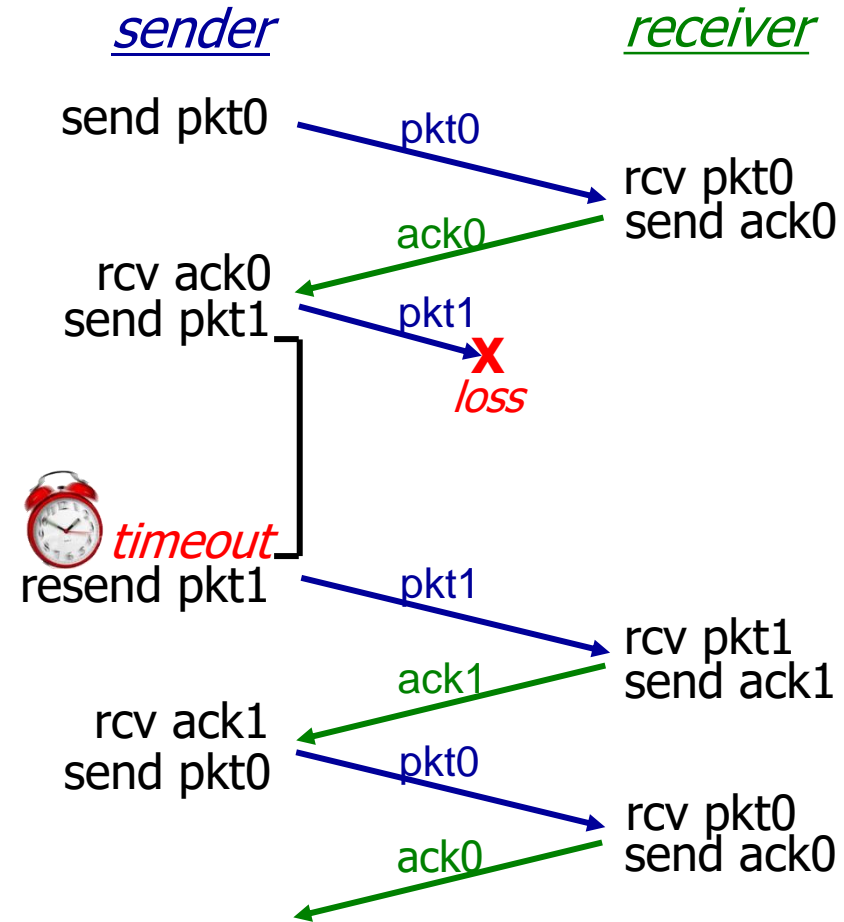
“reasonable” amount of time for ACK

- ❖ retransmits if no ACK received in this time
- ❖ if pkt (or ACK) just delayed (not lost):
  - retransmission will be duplicate, but use of seq. #'s already handles this
  - receiver must specify seq # of pkt being ACKed
- ❖ requires countdown timer

# rdt3.0 in action

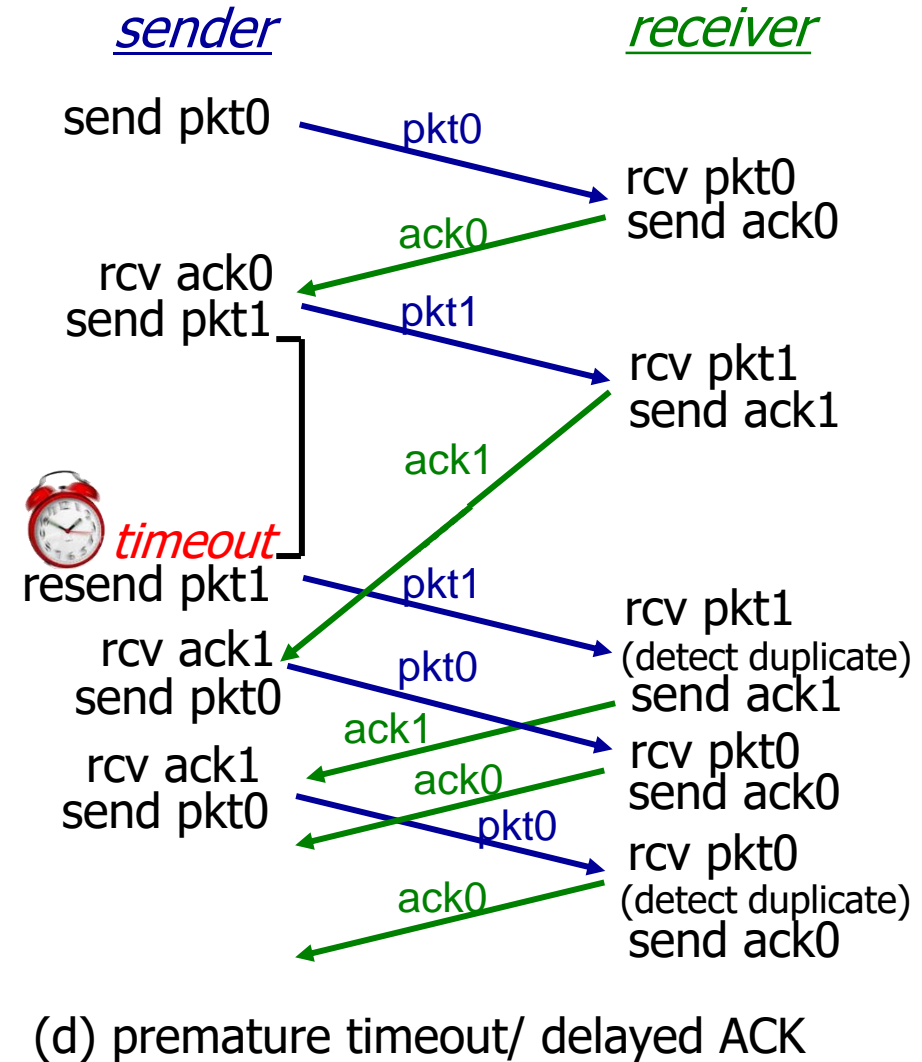
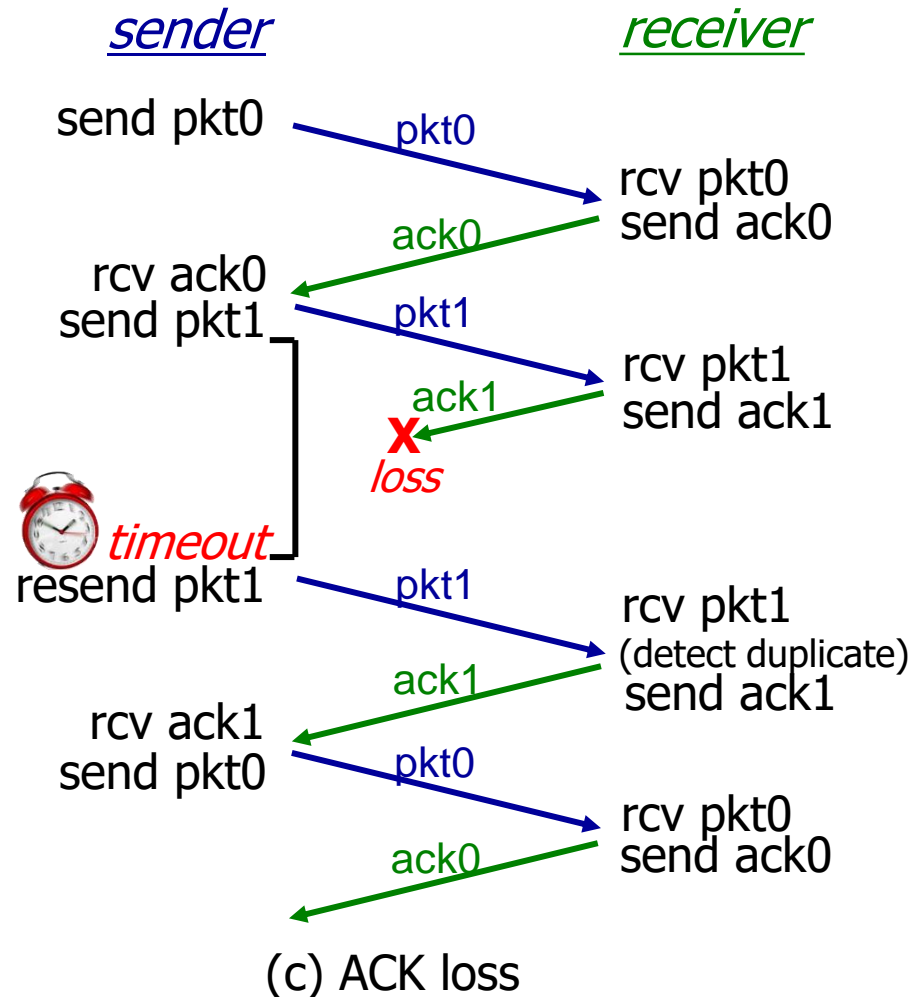


(a) no loss



(b) packet loss

# rdt3.0 in action



# Recap

Rdt Version	Scenario	Features Added
1.0	No error	Nothing
2.0	Data Bit Error	Checksum, ACK/NAK
2.1	Data Bit Error ACK/NAK Bit Error	Checksum, ACK/NAK, Sequence Number
3.0	Data Bit Error ACK/NAK Bit Error Packet Loss	Checksum, ACK/NAK, Sequence Number, Timeout



# Performance of rdt3.0

- ❖ rdt3.0 works, but performance stinks
- ❖ ex: 1 Gbps link, 15 ms prop. delay, 8000 bit packet:


$$d_{trans} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bps}} = 8 \text{ microseconds}$$

- $U_{\text{sender}}$ : **utilization** - fraction of time sender busy sending

$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

- if  $RTT=30$  msec, 1KB pkt every 30 msec  $\rightarrow$  33kB/sec thruput over 1 Gbps link
- network protocol limits use of physical resources!

# rdt3.0: stop-and-wait operation




$$U_{\text{sender}} = \frac{L / R}{RTT + L / R} = \frac{.008}{30.008} = 0.00027$$

Why send only one packet?  
Can we send more?

# Pipelined protocols

**pipelining:** sender allows multiple, "in-flight", yet-to-be-acknowledged pkts

- range of sequence numbers must be increased
- buffering at sender and/or receiver




(a) a stop-and-wait protocol in operation

(b) a pipelined protocol in operation

- ❖ two generic forms of pipelined protocols: **go-Back-N**, **selective repeat**

# Pipelining: increased utilization



$$U_{\text{sender}} = \frac{3 * L / R}{RTT + L / R} = \frac{.024}{30.008} = 0.0008$$

How many packets should be in the pipeline?

# Bandwidth-Delay Product (BDP)

- ❖ A logical link has a bandwidth of  $B$  bits/s and the round trip time is  $D$  sec.
- ❖ Bandwidth-Delay Product or BDP is computed as  $B \cdot D$
- ❖ Example:
  1. Two computers are connected over a WAN and can transfer data at 2Mbps with 100ms round trip time
    - $BDP = 2 \cdot 0.1 = 0.2 \text{ Mb}$
  2. A satellite link has a bandwidth of 10Mbps and RTT of 1s
    - $BDP = 10 \cdot 1 = 10 \text{ Mb}$
  3. Two computers from two universities in different continents are connected by a 500Mbps link and has RTT of 200ms
    - $BDP = 500 \cdot 0.2 = 100 \text{ Mb}$

# BDP Computation (cont'd)

How do you interpret BDP?  
Why is BDP interesting?

- ❖ BDP estimates on the amount of data that the communicate pipe can "store"
  - If the amount of unack data (window size) allowed is less than the BDP, sender will be idle at times, leading to lower efficiency


# Pipelined Protocols

- Selective Repeat: big pic
  - sender can have up to N unack'ed packets in pipeline
  - rcvr sends individual ack for each packet
  - sender maintains timer for each unacked packet
    - when timer expires, retransmit only unack'ed packet
- Go-back-N: big picture:
  - sender can have up to N unacked packets in pipeline
  - rcvr only sends
  - cumulative acks
    - doesn't ack packet if
  - there's a gap
  - sender has timer for oldest unacked packet
    - if timer expires, retransmit all unack'ed packets

# Go-Back-N

## Sender:

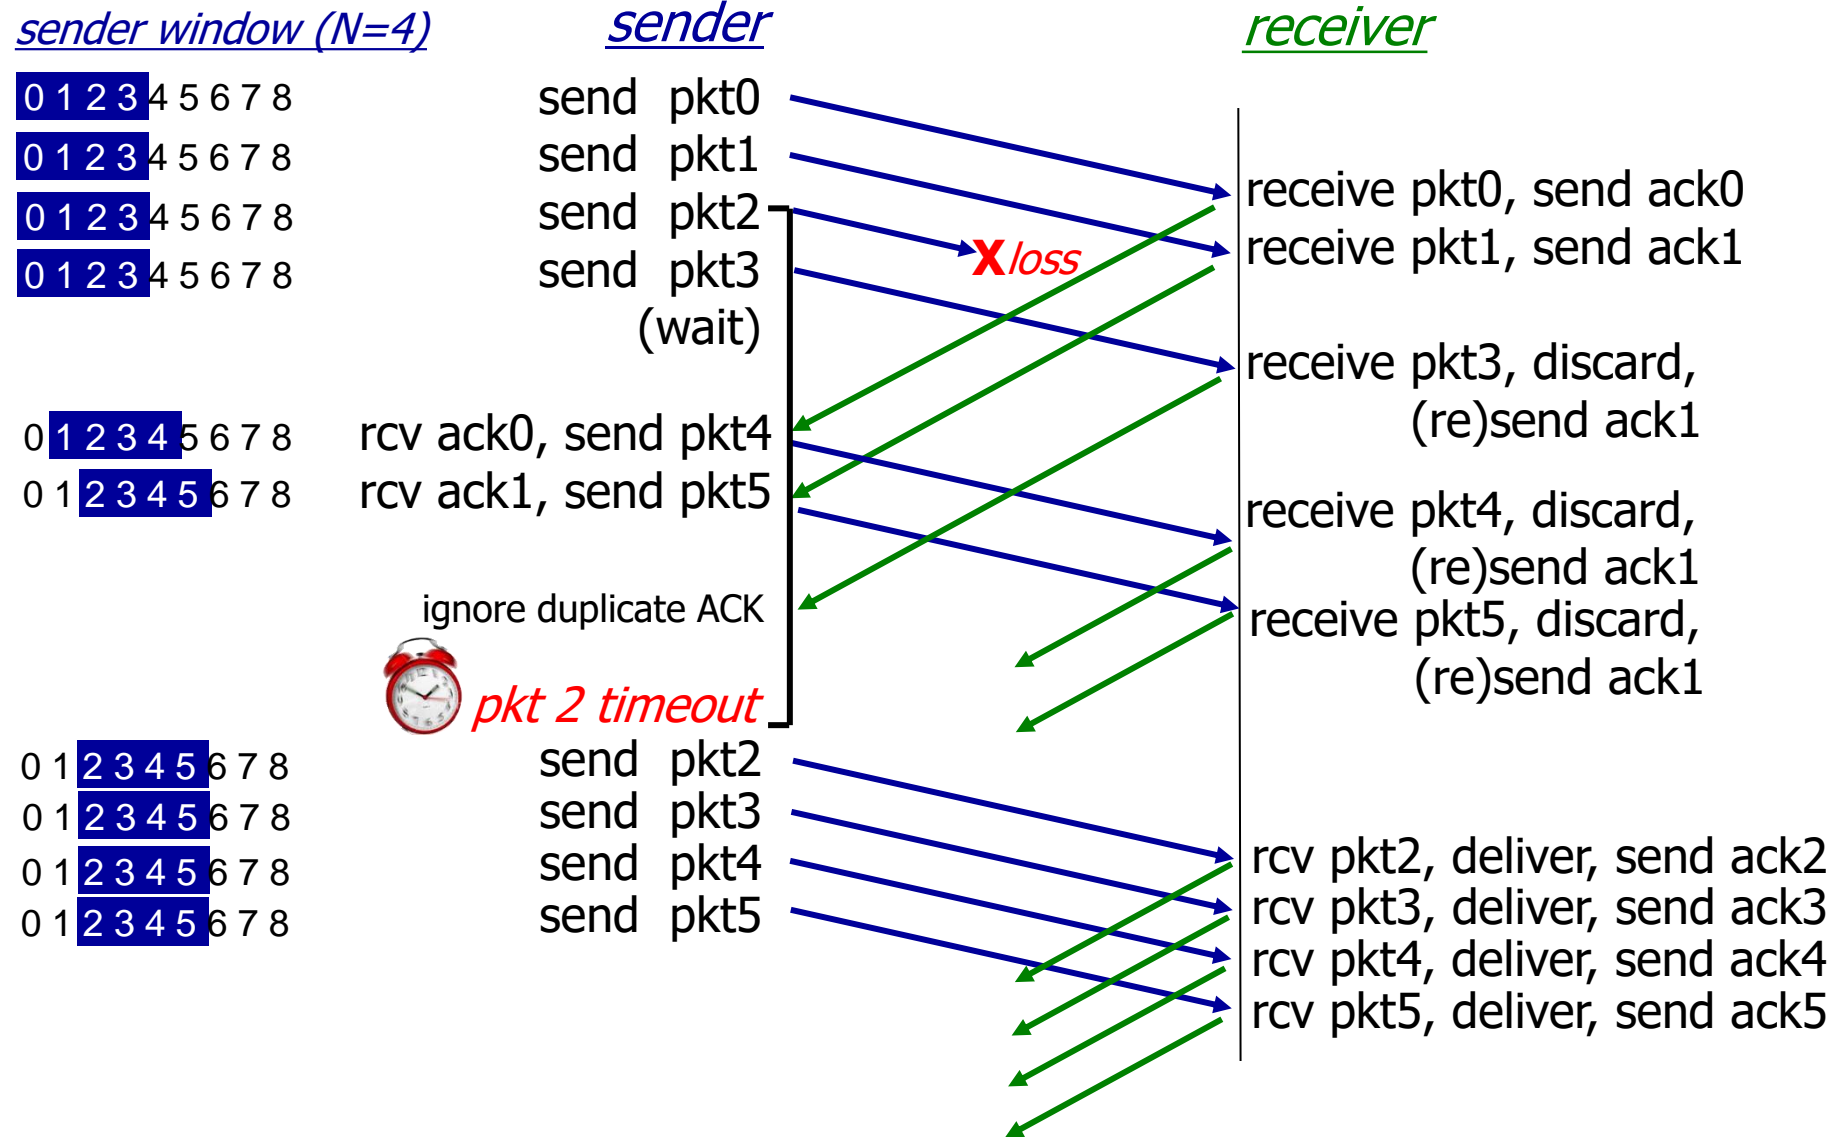
- ❖ k-bit seq # in pkt header
- ❖ "window" of up to N, consecutive unack'ed pkts allowed



- ❖ ACK(n): ACKs all pkts up to, including seq # n - "cumulative ACK"
  - may receive duplicate ACKs (see receiver)
- ❖ timer for each in-flight pkt
- ❖ timeout(n): retransmit pkt n and all higher seq # pkts in window



# Go-Back-N



# Selective Repeat

- ❖ receiver individually acknowledges all correctly received pkts
  - buffers pkts, as needed, for eventual in-order delivery to upper layer
- ❖ sender only resends pkts for which ACK not received
  - sender timer for each unACKed pkt
- ❖ sender window
  - N consecutive seq #'s
  - again limits seq #'s of sent, unACK'ed pkts